

Performance Analysis of Parallel K-Means

Color Quantization Optimization using Numba & Multi-Core
CPU

Pere Llauradó Adeva | Final Project

Problem Introduction

CONTEXT & MOTIVATION

Image Colors Quantization

- Reduce the number of colors while preserving visual fidelity.
- Essential for image compression.
- Fastest processing utile for example in medical or biological images



COMPUTATIONAL CHALLENGES



Complexity

Up to 6.0×10^9 operations for 1024x1024 pixels images using 32 clusters. This leads to massive CPU load during distance calculations.



The GIL & Interpreter

Python's Global Interpreter Lock prevents real multithreading. Standard loops cannot use multiple cores effectively.



Race Conditions

Centroid updates risk data corruption if multiple threads write to the same memory address simultaneously.

The Approach

K-MEANS ALGORITHM

1. Initialization (Pre-loop)

- ✓ Converts the image into a long list of pixels
- ✓ Select K random pixels from the original img. to act as the initial color centers

3. Final Reconstruction

- ✓ Once the centroids are finally determined
- ✓ Assign to each pixel the colour of its centroid

2. Update Phase

A- Assignment:

Calculate distance for every pixel to each of the centroids and assign a label to the pixel

B- Update:

Calculate new centroids by taking the mean of the color values of all pixels assigned to that cluster

C- Convergence:

- How much the centroids have moved
- If the movement is less than the tolerance the algorithm stops early.

PARALLELIZED PARTS

2.A. Assignment (Prange)

- ✓ Parallel pixel distribution via Numba prange.
- ✓ Each thread processes a chunk of pixels simultaneously

2.B. Update Phase (Manual)

- ⟳ Manual chunking to manage thread-local data.
- ⟳ Each thread calculates its own partial sums

2.C. Reconstruction

- ✓ Parallel pixel distribution via Numba prange.
- ✓ Each thread processes a chunk of pixels simultaneously

DATASETS USED

Dataset 1 – Landscape Pictures:

High resolution, RGB pictures with resolutions of: 1024x768,
1024x1024, 1600x319... pixels



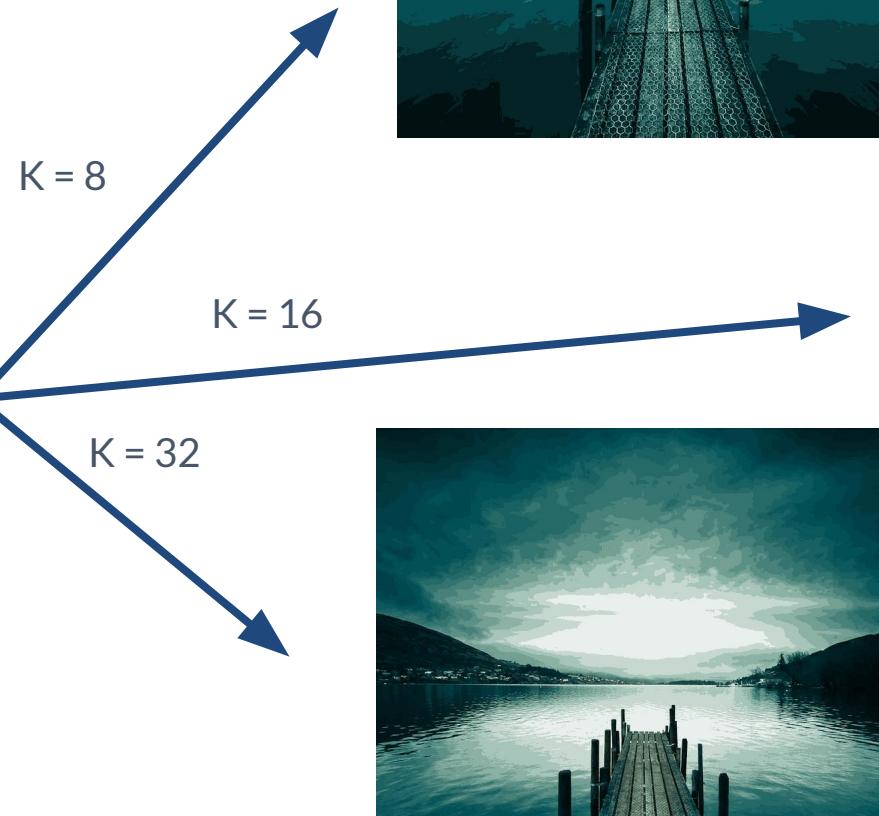
Dataset 2 – Intel Image Classification:

Low resolution, RGB pictures with a resolution of 150x150
pixels



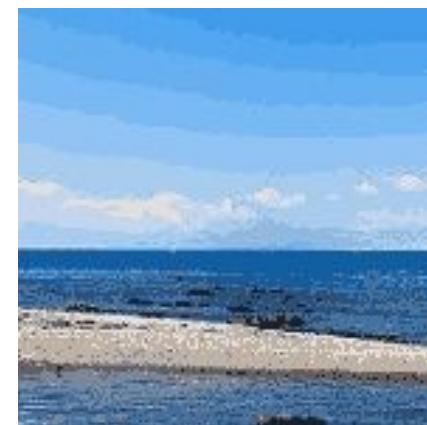
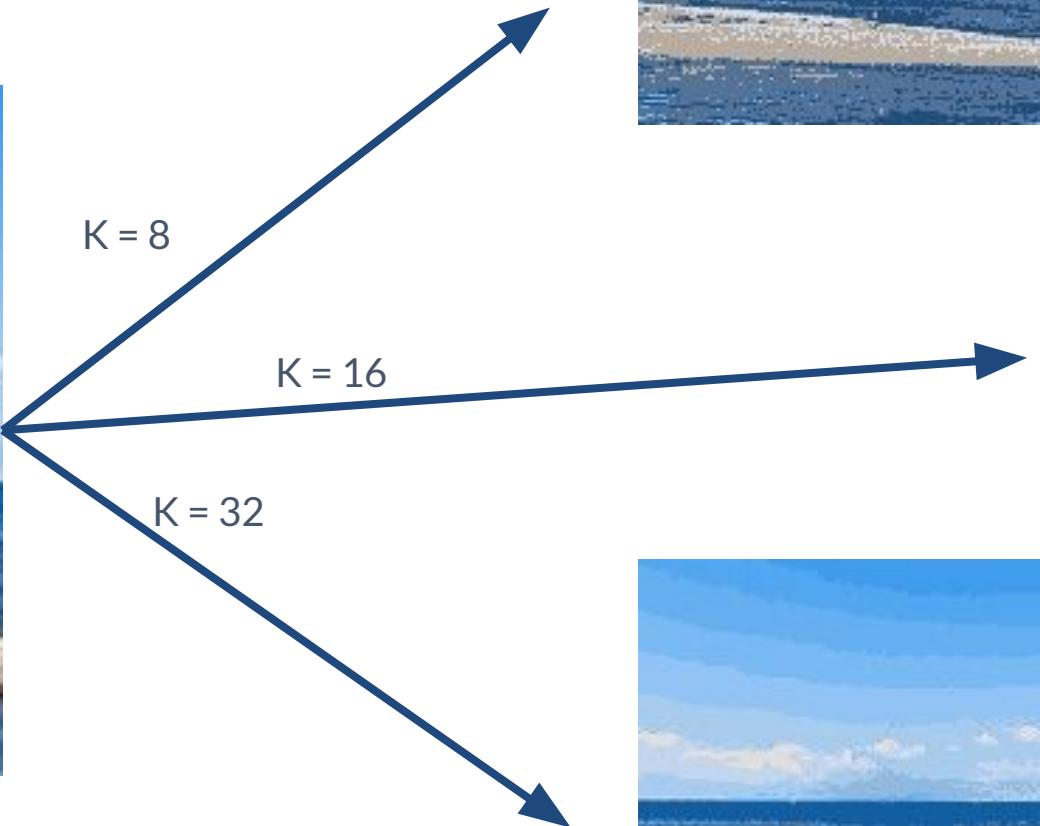
K-MEANS VISUAL RESULTS

Dataset 1 — Landscape Pictures:



K-MEANS VISUAL RESULTS

Dataset 2 – Intel Image Classification:



DATASETS & WORKLOAD COMPLEXITY

Dataset Name	Resolution	Pixels (N)	Ops (K=32)
Small_Res (Intel Classification)	150 x 150	22,500	$\sim 1.3 \times 10^8$
Full_Res (Landscape Pictures)	1024 x 1024	1,048,576	$\sim 6.0 \times 10^9$

PARALLELIZATION STRATEGY

Native Speed Execution

By using Numba's `@jit(parallel=True)`, we bypass the GIL entirely and compile the algorithm into LLVM machine code.

This allows the 8 logical threads of our Core i7 to saturate memory bandwidth and maximize floating-point throughput.

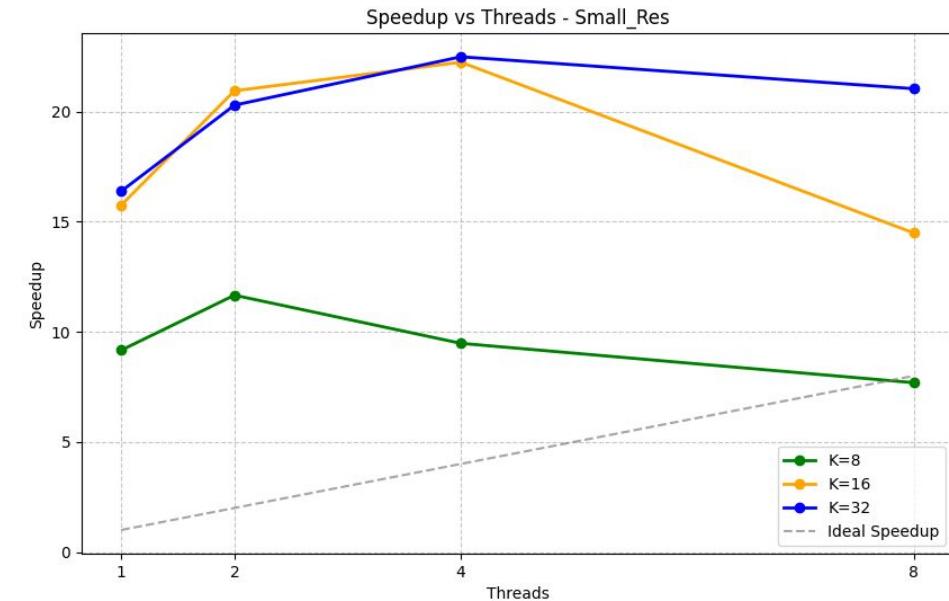
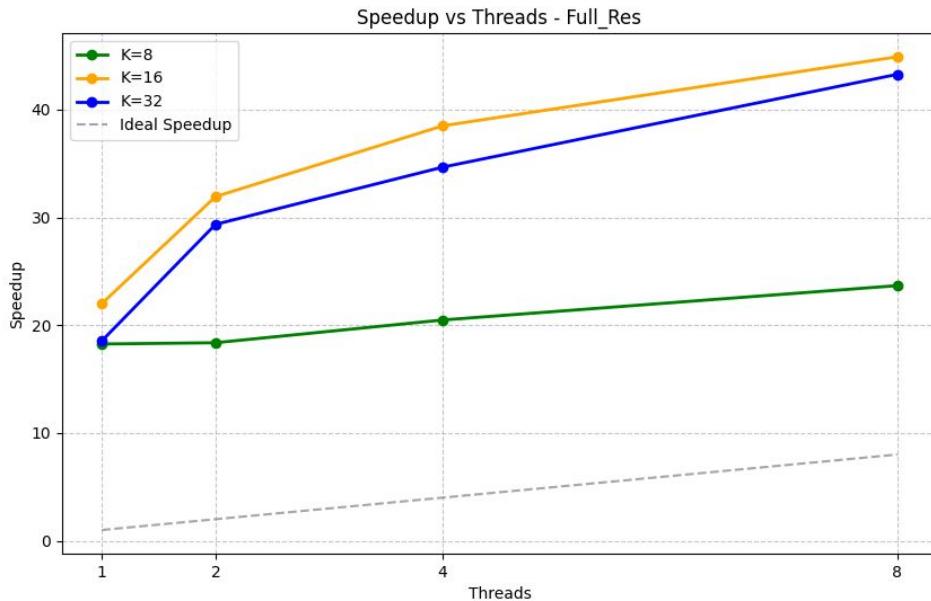
Automatic and manual work distribution: using orange when the operation is fully independent. In the risky cases we use manual distribution.

KEY OPTIMIZATIONS

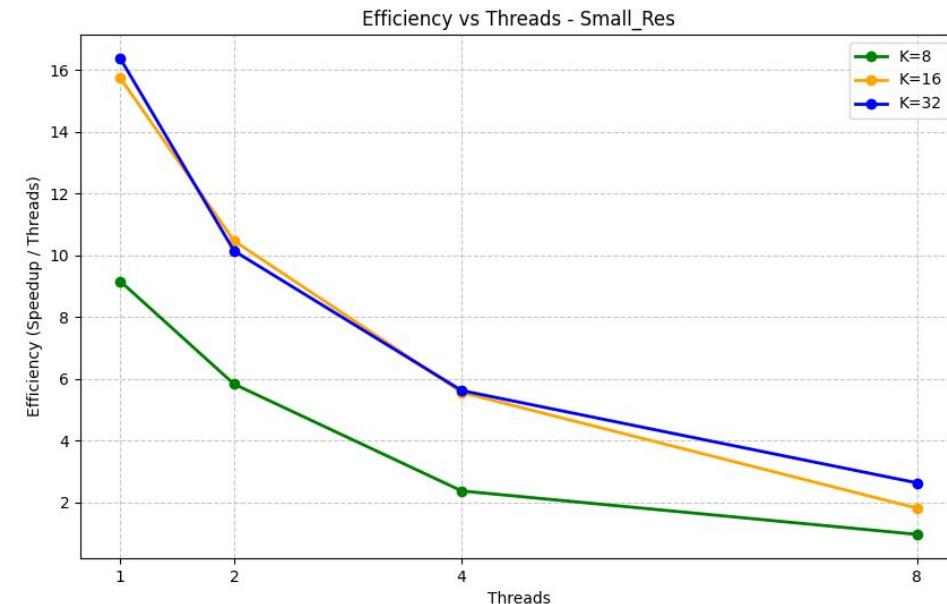
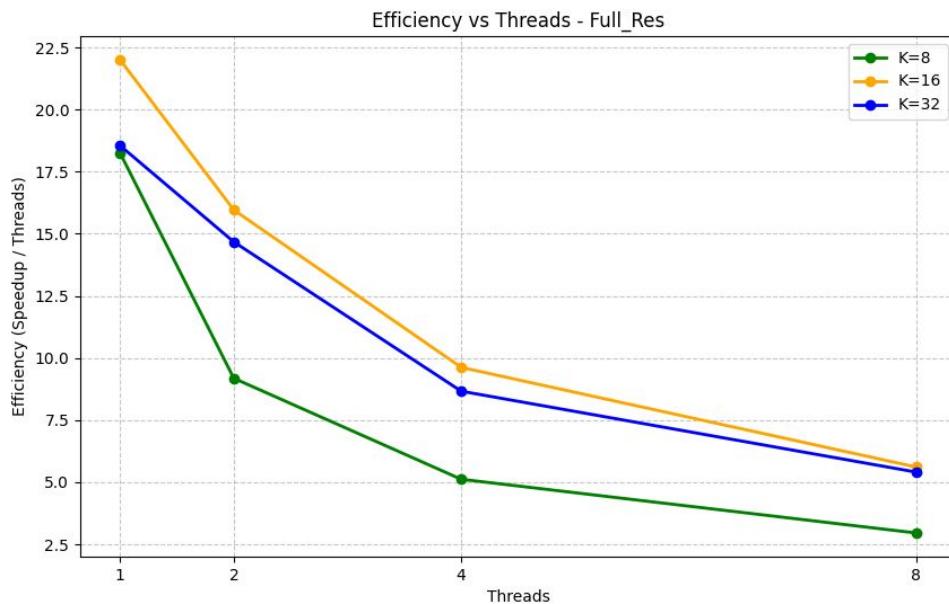
-  **False Sharing Prevention:** Added 64-byte padding to private thread buffers, ensuring data stays on separate L1 cache lines.
-  **JIT Warm-up:** Pre-compiling functions before measurement to exclude overhead from LLVM compilation time.
-  **Array Flattening:** Reshaping 3D images to 2D contiguous arrays for optimal CPU cache prefetching.
-  **Inertia Validation:** Comparing Sum of Squared Errors rather than exact colors to calculate the accuracy of the algorithm

Performance Results

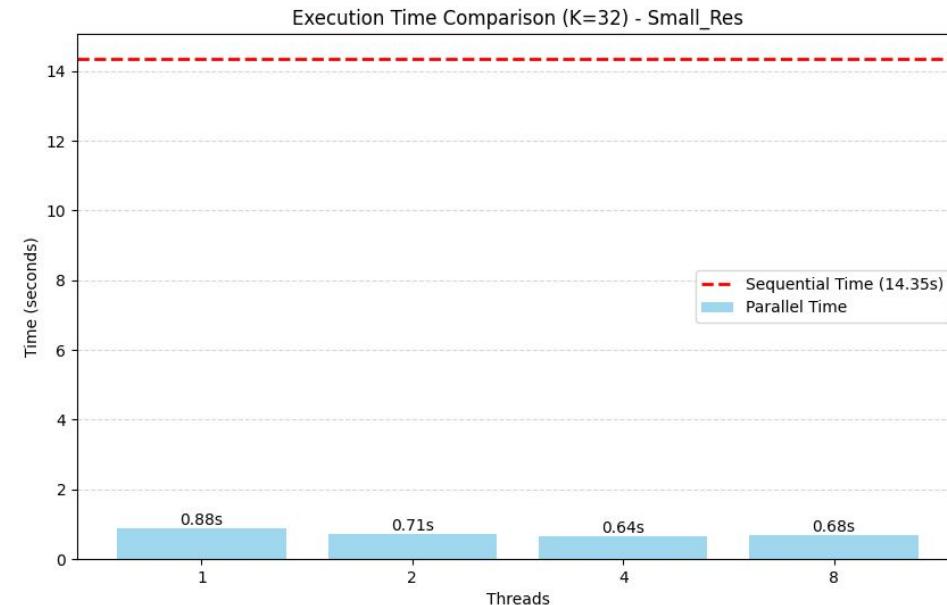
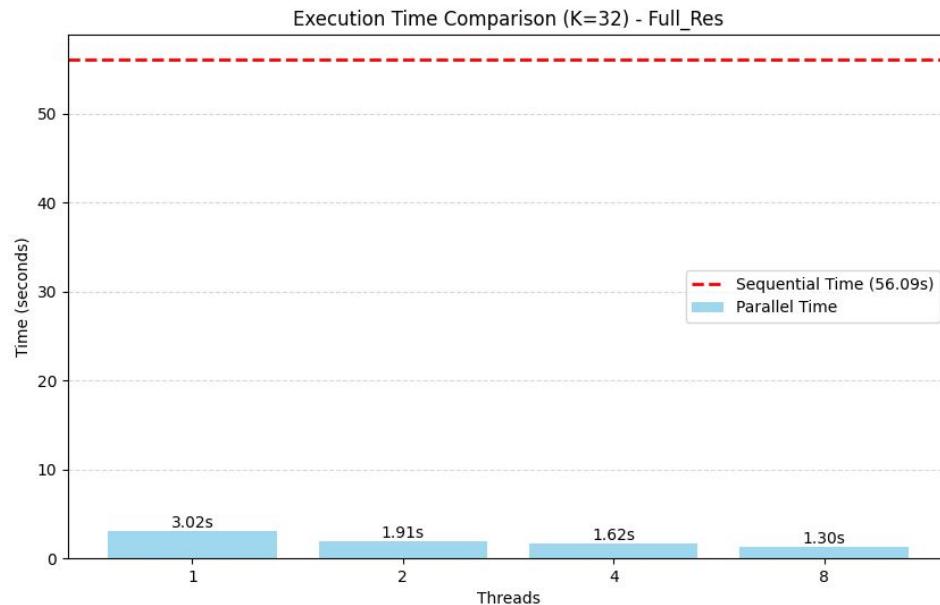
SPEEDUP VS THREADS



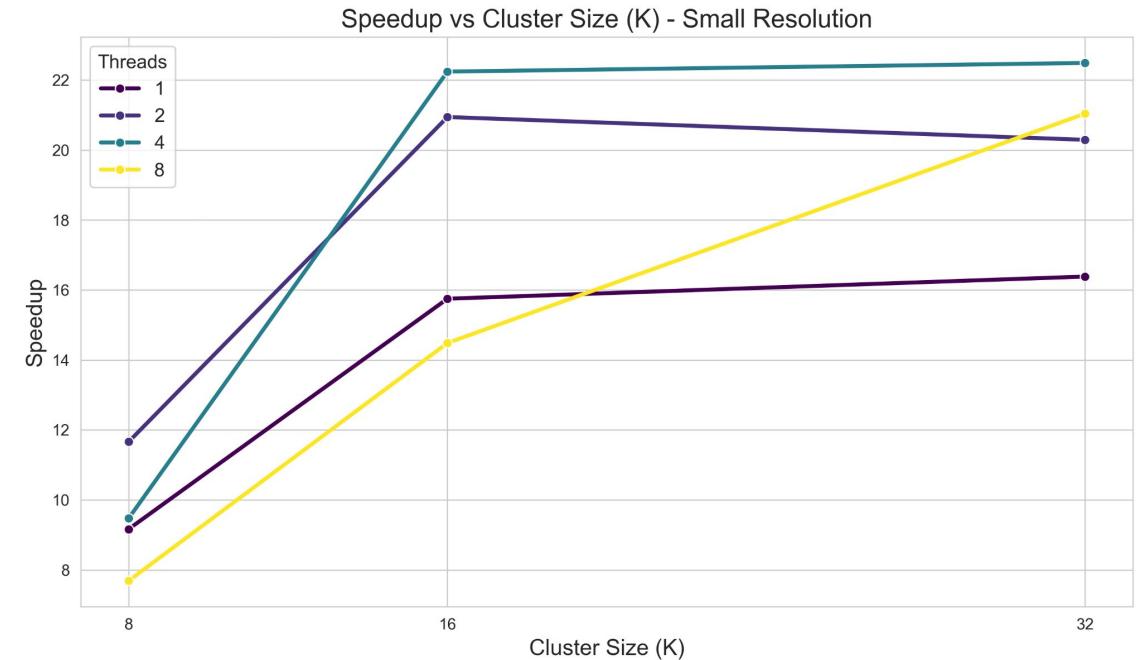
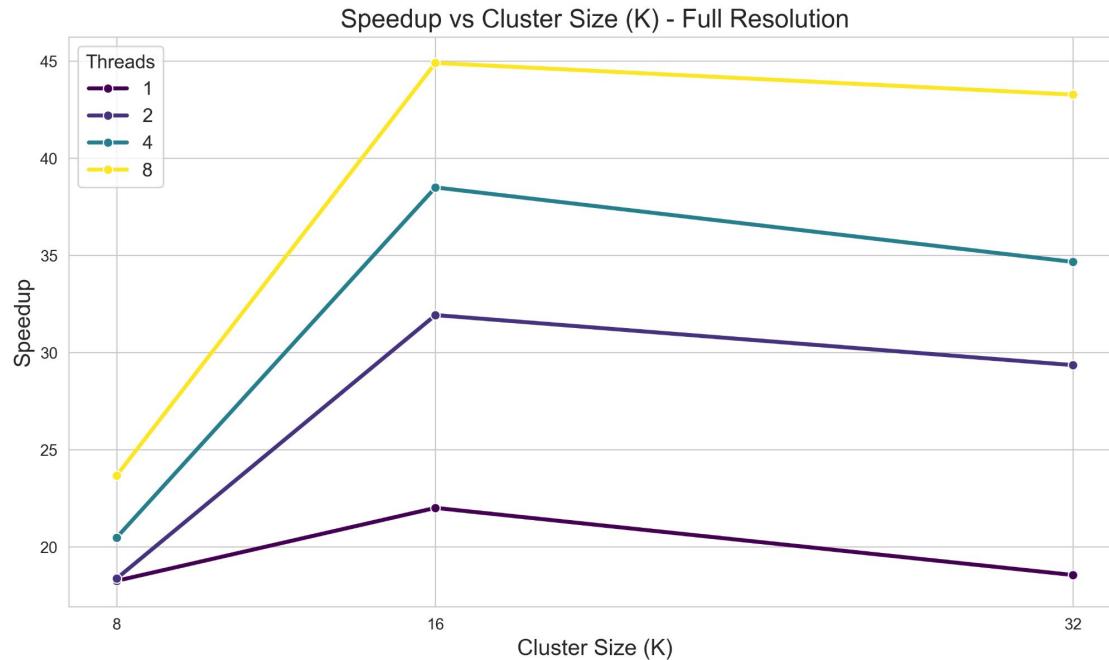
EFFICIENCY VS THREADS



TIME VS THREADS



SPEEDUP VS CLUSTER SIZE



BEST CONFIGURATIONS

Hw settings: i7-8550U with 4 physical cores and 8 logical processors utilizing Hyper-Threading

Threads: 8 is the best number of threads for Dataset 1 and 4 for Dataset 2

Clusters: for the Dataset 1 is K = 16 and for the Dataset 2 is K = 32

Dataset 1: High resolution

44.9x Speedup

Dataset 2: Low resolution

22.49x SpeedUp

INTERESTING FINDINGS OR SURPRISES

Validation problem

- Floating-Point sum error: float limited size make microscopic rounding differences
- The "Butterfly Effect" in K-Means: A tiny rounding error can flip a single pixel to a different cluster.
- Different correct solutions
- Total error comparison

Empty clusters

- On the small dataset, when we tested with $K = 32$, some clusters ends up being orphaned, without any pixel belonging to it.
- **First Idea:** pick a random pixel as the new centroid
- **Correct idea:** leave the centroid as black [0,0,0] == deleting it, to avoid parallelization problems

Conclusions & Lessons

SUMMARY

Performance gains

- Reduce the processing time from 56 s to 1.29 s with a 44,9x speedup.
- Numba can pass the python GIL to accelerate and parallelize heavy image processing

Optimal configuration

- **High resolution dataset:** 8 threads an K = 16
- **Low resolution dataset:** 4 threads and K = 32

Scalability

- **Full_Res images:** scale significantly better (Gustafson's Law) as the computational load hides thread management overhead.
- **Small_Res images:** are limited by Amdahl's Law, where fixed synchronization costs dominate the short execution time.

LESSONS LEARNED

- **Power of compilation:** eliminating the python interpreter reduces more the execution time than using threads
- **Cache sizes:** understand the size of the cache and add the necessary padding to avoid threads fighting for memory access.
- **Validation correctness importance:** depending on the algorithm and the size of the problem the result of the same execution could be slightly different. Important to consider it for the tests.

Thank You!

Questions or Discussion

Bibliografy

- [1] Dataset 1 (Big images):Kaggle - Landscape Pictures
- <https://www.kaggle.com/datasets/arnaud58/landscape-pictures?resource=download>
- [2] Dataset 2 (Small images):Kaggle - Intel Image Classification
- <https://www.kaggle.com/datasets/puneet6060/intel-image-classification?resource=download>
- [3] Parallel Programming Course Lecture Notes and Slides. Course material from the Parallel Programming course.